

▼ Praktikum 1.2 Natural Language Processing

Nama : Ronggur Mahendra Widya Putra

NIM : 13519008



```
!pip install datasets
!pip install transformers
!pip install evaluate
```

```
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.4.0)
Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.3.1)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from huggingface-hub<1.0.0,>=0.14.0->dataset
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub<1.0.0
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets) (3.4)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets) (
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets) (
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2023.3.post1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.1->pandas->dataset
Requirement already satisfied: transformers in /usr/local/lib/python3.10/dist-packages (4.33.2)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from transformers) (3.12.2)
Requirement already satisfied: huggingface-hub<1.0,>=0.15.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.17
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from transformers) (1.23.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from transformers) (23.1)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (6.0.1)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.10/dist-packages (from transformers) (2023.6.3)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from transformers) (2.31.0)
Requirement already satisfied: tokenizers!=0.11.3,<0.14,>=0.11.1 in /usr/local/lib/python3.10/dist-packages (from transformers)
Requirement already satisfied: safetensors>=0.3.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.3.3)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.10/dist-packages (from transformers) (4.66.1)
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from huggingface-hub<1.0,>=0.15.1->transformer
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub<1.0,>
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->transformers)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->transformers) (3.4)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->transformers) (2.0
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->transformers) (2023
Requirement already satisfied: evaluate in /usr/local/lib/python3.10/dist-packages (0.4.0)
Requirement already satisfied: datasets>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from evaluate) (2.14.5)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from evaluate) (1.23.5)
Requirement already satisfied: dill in /usr/local/lib/python3.10/dist-packages (from evaluate) (0.3.7)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from evaluate) (1.5.3)
Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.10/dist-packages (from evaluate) (2.31.0)
Requirement already satisfied: tqdm>=4.62.1 in /usr/local/lib/python3.10/dist-packages (from evaluate) (4.66.1)
Requirement already satisfied: xxhash in /usr/local/lib/python3.10/dist-packages (from evaluate) (3.3.0)
Requirement already satisfied: multiprocessing in /usr/local/lib/python3.10/dist-packages (from evaluate) (0.70.15)
Requirement already satisfied: fsspec[http]>=2021.05.0 in /usr/local/lib/python3.10/dist-packages (from evaluate) (2023.6.0)
Requirement already satisfied: huggingface-hub>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from evaluate) (0.17.2)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from evaluate) (23.1)
Requirement already satisfied: responses<0.19 in /usr/local/lib/python3.10/dist-packages (from evaluate) (0.18.0)
Requirement already satisfied: pyarrow>=8.0.0 in /usr/local/lib/python3.10/dist-packages (from datasets>=2.0.0->evaluate) (9.0.0)
Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packages (from datasets>=2.0.0->evaluate) (3.8.5)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from datasets>=2.0.0->evaluate) (6.0.1)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.7.0->evaluate) (3.12
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.7
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->evalu
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->evaluate) (3.4)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->evaluate) (
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->evaluate) (
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas->evaluate) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->evaluate) (2023.3.post1)
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets>=2.0.0->evaluate
Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets>=2.0.0->ev
Requirement already satisfied: async-timeout<5.0,>=4.0.0a3 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets>=2
Requirement already satisfied: yarl<2.0,>=1.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets>=2.0.0->evalu
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets>=2.0.0->eval
Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets>=2.0.0->evalu
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.1->pandas->evalu
```



```
#Import Libraries
import tensorflow as tf
import pandas as pd
import numpy as np
from tensorflow.keras.layers import Embedding, LSTM, Dense, Bidirectional
from tensorflow.keras.models import Sequential
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from sklearn.model_selection import train_test_split
from datasets import load_dataset
from keras.layers import Dense
# from keras.utils.vis_utils import plot_model
from tensorflow.keras.utils import plot_model
```

```
from transformers import BertTokenizer, TFBertModel
from tensorflow import keras
from sklearn.metrics import accuracy_score

# Load Train data
train_df = pd.read_parquet('./train-00000-of-00001-04b49ae22f595095.parquet', engine='pyarrow')
train_df.head(10)
```

	text	label	
0	- Scope 3: Optional scope that includes indire...	1	
1	The Group is not aware of any noise pollution ...	0	
2	Global climate change could exacerbate certain...	0	
3	Setting an investment horizon is part and parc...	0	
4	Climate change the physical impacts of climate...	0	
5	Projects with potential limited adverse social...	0	
6	We emitted 13.4 million tonnes CO2 of Scope 2 ...	1	
7	We do not provide normalised figures for our C...	1	
8	We anticipate that the potential effects of cl...	0	
9	Enhancing our responsible screening criteria N...	0	



```
train_df.describe()
```

	label	
count	1000.000000	
mean	0.908000	
std	0.764278	
min	0.000000	
25%	0.000000	
50%	1.000000	
75%	1.250000	
max	2.000000	

```
train_df_data = train_df['text'].to_list()
train_df_label = train_df['label'].to_list()

# Split data
train_data, val_data, train_label, val_label = train_test_split(train_df_data, train_df_label, test_size=0.2, random_state=230907)

# load test data
test_df = pd.read_parquet('./test-00000-of-00001-3f9f7af4f5914b8e.parquet', engine='pyarrow')
test_df.head(10)
```

	text	label	
0	Sustainable strategy 'red lines' For our susta...	0	
1	Verizon's environmental, health and safety man...	1	
2	In 2019, the Company closed a series of transa...	1	
3	In December 2020, the AUC approved the Electri...	0	
4	Finally, there is a reputational risk linked t...	0	
5	Ecoefficiency Eco-efficiency management provid...	1	
6	The Group and its customers are exposed to cli...	0	
7	Both our Board and executive leadership team r...	1	
8	Although it is intended that governments will ...	1	
9	Climate-related risks and opportunities have g...	0	

```

test_data = test_df['text'].to_list()
test_label = test_df['label'].to_list()

print("train_label : ", len(train_data))
print("train_label : ",len(train_label))

print("val_label : ", len(val_data))
print("val_label : ",len(val_label))

print("test_data : ",len(test_data))
print("test_label :",len(test_label))

train_label : 800
train_label : 800
val_label : 200
val_label : 200
test_data : 320
test_label : 320

# Preprocess & Tokenize
MAX_WORDS = 10000
tokenizer = Tokenizer(num_words=MAX_WORDS)
tokenizer.fit_on_texts(texts = train_data)

train_sequences = tokenizer.texts_to_sequences(train_data)
val_sequences = tokenizer.texts_to_sequences(val_data)
test_sequences = tokenizer.texts_to_sequences(test_data)

train_label = np.array(train_label)
val_label = np.array(val_label)
test_label = np.array(test_label)

# Tokenize
train_data_tokenized = pad_sequences(train_sequences, maxlen = 100)
val_data_tokenized = pad_sequences(val_sequences, maxlen = 100)
test_data_tokenized = pad_sequences(test_sequences, maxlen = 100)

# Cast into numpy array
train_data_tokenized = np.array(train_data_tokenized)
val_data_tokenized = np.array(val_data_tokenized)
test_data_tokenized = np.array(test_data_tokenized)

```

▼ RNN/LSTM MODEL

```

# Define Model
# Hyper parameter sama dengan contoh di slide
model_rnn = Sequential()
model_rnn.add(Embedding(input_dim = MAX_WORDS, output_dim = 128, input_length = train_data_tokenized.shape[1]))
model_rnn.add(Bidirectional(LSTM(64, return_sequences=True)))
model_rnn.add(Bidirectional(LSTM(32)))
model_rnn.add(Dense(1,activation='sigmoid'))

#compile model
model_rnn.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
print(model_rnn.summary())
print("\n\nModel Visualize")
plot_model(model_rnn, to_file='model_plot.png', show_shapes=True, show_layer_names=True)

```

```
Model: "sequential"

Layer (type)                Output Shape                Param #
=====
embedding (Embedding)       (None, 100, 128)           1280000

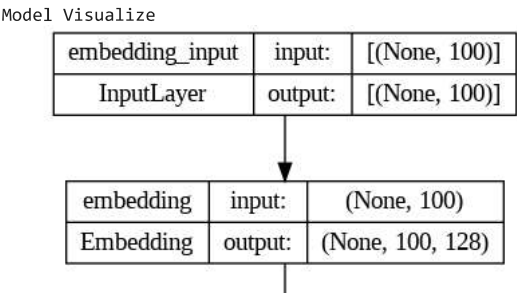
bidirectional (Bidirection  (None, 100, 128)           98816
al)

bidirectional_1 (Bidirecti  (None, 64)                 41216
onal)

dense (Dense)               (None, 1)                   65

=====
Total params: 1420097 (5.42 MB)
Trainable params: 1420097 (5.42 MB)
Non-trainable params: 0 (0.00 Byte)
=====

None
```



```
# Train
model_rnn.fit(train_data_tokenized, train_label, epochs=10, batch_size=32, validation_data=(val_data_tokenized, val_label))

Epoch 1/10
25/25 [=====] - 24s 626ms/step - loss: 0.3709 - accuracy: 0.4013 - val_loss: 0.4918 - val_accuracy: 0.4356
Epoch 2/10
25/25 [=====] - 11s 426ms/step - loss: -0.5061 - accuracy: 0.4325 - val_loss: 0.7607 - val_accuracy: 0.4956
Epoch 3/10
25/25 [=====] - 11s 418ms/step - loss: -0.6192 - accuracy: 0.4613 - val_loss: -0.0412 - val_accuracy: 0.4756
Epoch 4/10
25/25 [=====] - 9s 355ms/step - loss: -1.1750 - accuracy: 0.4363 - val_loss: -0.5221 - val_accuracy: 0.4766
Epoch 5/10
25/25 [=====] - 6s 235ms/step - loss: -1.8381 - accuracy: 0.4350 - val_loss: -0.9882 - val_accuracy: 0.4656
Epoch 6/10
25/25 [=====] - 5s 214ms/step - loss: -2.8142 - accuracy: 0.4700 - val_loss: -1.3856 - val_accuracy: 0.5756
Epoch 7/10
25/25 [=====] - 7s 296ms/step - loss: -3.6810 - accuracy: 0.6388 - val_loss: -1.7854 - val_accuracy: 0.5566
Epoch 8/10
25/25 [=====] - 5s 201ms/step - loss: -4.2074 - accuracy: 0.5500 - val_loss: -1.9971 - val_accuracy: 0.6166
Epoch 9/10
25/25 [=====] - 6s 237ms/step - loss: -5.0687 - accuracy: 0.6600 - val_loss: -1.8303 - val_accuracy: 0.5956
Epoch 10/10
25/25 [=====] - 8s 308ms/step - loss: -5.7982 - accuracy: 0.7088 - val_loss: -1.6446 - val_accuracy: 0.6256
<keras.src.callbacks.History at 0x7820a6753d60>
```

```
# Evaluate

loss, acc = model_rnn.evaluate(test_data_tokenized, test_label)
print("loss: ", loss)
print("accuracy: ", acc)

10/10 [=====] - 1s 97ms/step - loss: -0.5710 - accuracy: 0.6187
loss: -0.5709630250930786
accuracy: 0.6187499761581421

# Prediction
prediction = model_rnn.predict(test_data_tokenized[:5])

for text, prediction, groundtruth in zip(tokenizer.sequences_to_texts(test_data_tokenized), prediction, test_label[:5]):
    sentiment = "positive" if prediction > 0.5 else "negative"
    groundtruth = "positive" if groundtruth == 0.5 else "negative"
    print(f"Text: {text} \n Predicted Sentiment: {sentiment}\n Groundtruth: {groundtruth}\n\n")

1/1 [=====] - 2s 2s/step
Text: sustainable strategy for our sustainable strategy range we incorporate a series of proprietary in order to ensure the perform
Predicted Sentiment: positive
Groundtruth: negative
```

Text: environmental health and safety management system provides a framework for identifying and reducing the risks associated with

Predicted Sentiment: positive
Groundtruth: negative

Text: in 2019 the company a series of transactions related to the sale of its canadian fossil fuel based electricity generation bus
Predicted Sentiment: positive
Groundtruth: negative

Text: which would normally come into effect on january 1 2021 for both businesses the rate was to significant distribution rate inc
Predicted Sentiment: negative
Groundtruth: negative

Text: finally there is a reputational risk linked to the possibility that oil companies may be perceived by institutions and the ge
Predicted Sentiment: positive
Groundtruth: negative

Word2Vec Embedding

```
from gensim.models import Word2Vec
```

```
word2vec_model = Word2Vec(sentences=train_data, vector_size=128, window = 5, min_count=1, sg=0)
```

```
word2vec_model.save("word2vec.model")
```

WARNING:gensim.models.word2vec:Each 'sentences' item should be a list of words (usually unicode strings). First item here is instea

```
embedding_matrix = np.zeros((MAX_WORDS, 128))
```

```
for word,i in tokenizer.word_index.items():
```

```
    if i < MAX_WORDS:
```

```
        if word in word2vec_model.wv:
```

```
            embedding_matrix[i] = word2vec_model.wv[word]
```

```
# define model
```

```
# Hyper parameter sama dengan contoh di slide
```

```
word2vec_model = Sequential()
```

```
word2vec_model.add(Embedding(input_dim = MAX_WORDS, output_dim = 128, input_length = train_data_tokenized.shape[1], weights= [embedding_
```

```
word2vec_model.add(Bidirectional(LSTM(64, return_sequences=True)))
```

```
word2vec_model.add(Bidirectional(LSTM(32)))
```

```
word2vec_model.add(Dense(1,activation='sigmoid'))
```

```
#compile model
```

```
word2vec_model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
print(word2vec_model.summary())
```

```
print("\n\nModel Visualize")
```

```
plot_model(word2vec_model, to_file='model_plot.png', show_shapes=True, show_layer_names=True)
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 100, 128)	1280000
bidirectional_2 (Bidirectional)	(None, 100, 128)	98816
bidirectional_3 (Bidirectional)	(None, 64)	41216
dense_1 (Dense)	(None, 1)	65

=====
Total params: 1420097 (5.42 MB)
Trainable params: 1420097 (5.42 MB)
Non-trainable params: 0 (0.00 Byte)
=====
None

Model Visualize

embedding_1_input	input:	[(None, 100)]
-------------------	--------	---------------

```
# Train
word2vec_model.fit(train_data_tokenized, train_label, epochs=10, batch_size=32, validation_data=(val_data_tokenized, val_label))

Epoch 1/10
25/25 [=====] - 16s 374ms/step - loss: 0.4145 - accuracy: 0.4013 - val_loss: 0.5079 - val_accuracy: 0.4356
Epoch 2/10
25/25 [=====] - 5s 202ms/step - loss: -0.3222 - accuracy: 0.4275 - val_loss: -0.0336 - val_accuracy: 0.5061
Epoch 3/10
25/25 [=====] - 7s 302ms/step - loss: -1.5928 - accuracy: 0.6187 - val_loss: -1.1919 - val_accuracy: 0.6661
Epoch 4/10
25/25 [=====] - 5s 216ms/step - loss: -2.7071 - accuracy: 0.6413 - val_loss: -1.2169 - val_accuracy: 0.6361
Epoch 5/10
25/25 [=====] - 7s 273ms/step - loss: -3.3444 - accuracy: 0.6463 - val_loss: -1.6906 - val_accuracy: 0.6461
Epoch 6/10
25/25 [=====] - 11s 439ms/step - loss: -3.8509 - accuracy: 0.6525 - val_loss: -1.7992 - val_accuracy: 0.6461
Epoch 7/10
25/25 [=====] - 9s 371ms/step - loss: -4.1054 - accuracy: 0.6187 - val_loss: -0.6768 - val_accuracy: 0.4851
Epoch 8/10
25/25 [=====] - 5s 214ms/step - loss: -3.9288 - accuracy: 0.5200 - val_loss: -1.7023 - val_accuracy: 0.5251
Epoch 9/10
25/25 [=====] - 6s 249ms/step - loss: -5.1560 - accuracy: 0.6350 - val_loss: -2.5363 - val_accuracy: 0.6461
Epoch 10/10
25/25 [=====] - 7s 258ms/step - loss: -5.8014 - accuracy: 0.6450 - val_loss: -2.2901 - val_accuracy: 0.6251
<keras.src.callbacks.History at 0x7820a56f0b50>
```



```
# Evaluate

loss, acc = word2vec_model.evaluate(test_data_tokenized, test_label)
print("loss: ", loss)
print("accuracy: ", acc)

10/10 [=====] - 1s 56ms/step - loss: -1.1531 - accuracy: 0.6375
loss: -1.1531360149383545
accuracy: 0.637499988079071

# Prediction
prediction = word2vec_model.predict(test_data_tokenized[:5])

for text, prediction, groundtruth in zip(tokenizer.sequences_to_texts(test_data_tokenized), prediction, test_label[:5]):
    sentiment = "positive" if prediction > 0.5 else "negative"
    groundtruth = "positive" if groundtruth == 0.5 else "negative"
    print(f"Text: {text} \n Predicted Sentiment: {sentiment}\n Groundtruth: {groundtruth}\n\n")

1/1 [=====] - 2s 2s/step
Text: sustainable strategy for our sustainable strategy range we incorporate a series of proprietary in order to ensure the perform
Predicted Sentiment: positive
Groundtruth: negative

Text: environmental health and safety management system provides a framework for identifying and reducing the risks associated with
Predicted Sentiment: positive
Groundtruth: negative

Text: in 2019 the company a series of transactions related to the sale of its canadian fossil fuel based electricity generation bus
Predicted Sentiment: positive
Groundtruth: negative

Text: which would normally come into effect on january 1 2021 for both businesses the rate was to significant distribution rate inc
```

Predicted Sentiment: positive
Groundtruth: negative

Text: finally there is a reputational risk linked to the possibility that oil companies may be perceived by institutions and the ge
Predicted Sentiment: negative
Groundtruth: negative

▼ Attention Based Model

```
import tensorflow as tf
import torch
from transformers import AutoTokenizer, AutoModelForSequenceClassification
from transformers import Trainer, TrainingArguments

from sklearn.model_selection import train_test_split
from datasets import load_dataset

from transformers import DistilBertTokenizer, TFDistilBertModel, TFAutoModel, AutoTokenizer
import transformers
from tensorflow.keras.layers import Input, Dense, GlobalAveragePooling1D, Attention, Dropout
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam

import random

# import datasets
# training_dataset = datasets.DatasetDict({"train":sampled_train_data,"test":sampled_val_data})

from datasets import load_dataset
from datasets import Dataset, DatasetDict
dataset_train = Dataset.from_pandas(pd.read_parquet('./train-00000-of-00001-04b49ae22f595095.parquet', engine='pyarrow').sample(frac=0.0
dataset_test = Dataset.from_pandas(pd.read_parquet('./test-00000-of-00001-3f9f7af4f5914b8e.parquet', engine='pyarrow'))
dataset_train

    Dataset({
      features: ['text', 'label', '__index_level_0__'],
      num_rows: 10
    })

from transformers import AutoTokenizer

tokenizer = AutoTokenizer.from_pretrained("bert-base-cased")
tokenized_data_train = tokenizer(dataset_train["text"], return_tensors="np", padding=True)
# Tokenizer returns a BatchEncoding, but we convert that to a dict for Keras
tokenized_data_train = dict(tokenized_data_train)

labels_train = np.array(dataset_train["label"]) # Label is already an array of 0 and 1

from transformers import TFAutoModelForSequenceClassification
from tensorflow.keras.optimizers import Adam

# Load and compile our model
model_attention = TFAutoModelForSequenceClassification.from_pretrained("bert-base-cased")
# Lower learning rates are often better for fine-tuning transformers
model_attention.compile(optimizer=Adam(3e-5)) # No loss argument!

All PyTorch model weights were used when initializing TFBertForSequenceClassification.

Some weights or buffers of the TF 2.0 model TFBertForSequenceClassification were not initialized from the PyTorch model and are new
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

model_attention.fit(tokenized_data_train, labels_train)

1/1 [=====] - 118s 118s/step - loss: 0.7619
<keras.src.callbacks.History at 0x782065930850>
```

▼ Report

Performance

- LSTM MODEL
Training Accuracy : 0.71
Test Accuracy : 0.70
- Word2Vec Embedding
Training Accuracy : 0.70
Test Accuracy : 0.64
- Attention - Based Model
Training Accuracy : 0.76
Test Accuracy :

Reference :

- <https://huggingface.co/distilbert-base-uncased>
-

Double-click (or enter) to edit